

Task 2

February 7, 2018

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()

%matplotlib inline

In [2]: from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
from sklearn.neighbors import NearestCentroid
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV, KFold
from sklearn.preprocessing import StandardScaler

In [3]: from sklearn.datasets import fetch_covtype

In [4]: data = fetch_covtype(data_home='task2_data')

In [5]: data

Out[5]: {'DESCR': 'Forest covertype dataset.\n\nA classic dataset for classification benchmarks,
  'data': array([[ 2.59600000e+03,  5.10000000e+01,  3.00000000e+00, ...,
    0.00000000e+00,  0.00000000e+00,  0.00000000e+00],
  [ 2.59000000e+03,  5.60000000e+01,  2.00000000e+00, ...,
    0.00000000e+00,  0.00000000e+00,  0.00000000e+00],
  [ 2.80400000e+03,  1.39000000e+02,  9.00000000e+00, ...,
    0.00000000e+00,  0.00000000e+00,  0.00000000e+00],
  ...,
  [ 2.38600000e+03,  1.59000000e+02,  1.70000000e+01, ...,
    0.00000000e+00,  0.00000000e+00,  0.00000000e+00],
  [ 2.38400000e+03,  1.70000000e+02,  1.50000000e+01, ...,
    0.00000000e+00,  0.00000000e+00,  0.00000000e+00],
  [ 2.38300000e+03,  1.65000000e+02,  1.30000000e+01, ...,
    0.00000000e+00,  0.00000000e+00,  0.00000000e+00]]),
  'target': array([5, 5, 2, ..., 3, 3, 3], dtype=int32)}
```

```
In [6]: cov = pd.DataFrame(data['data'])
cov['Cover_Type'] = data['target']
```

```

In [7]: # Get column names
feature_file = open('feature_names','r')
features = []
for line in feature_file:
    features.append(line.split()[0])

In [8]: wildareas = [features[10] + str(i) for i in range(1,5)]
soiltypes = [features[11] + str(i) for i in range(1,41)]
features = features[:10] + wildareas + soiltypes + [features[12]]

In [9]: cov.columns = features

In [10]: cov.head()

```

	Elevation	Aspect	Slope	Horizontal_Distance_To_Hydrology	\
0	2596.0	51.0	3.0	258.0	
1	2590.0	56.0	2.0	212.0	
2	2804.0	139.0	9.0	268.0	
3	2785.0	155.0	18.0	242.0	
4	2595.0	45.0	2.0	153.0	

	Vertical_Distance_To_Hydrology	Horizontal_Distance_To_Roadways	\
0	0.0	510.0	
1	-6.0	390.0	
2	65.0	3180.0	
3	118.0	3090.0	
4	-1.0	391.0	

	Hillshade_9am	Hillshade_Noon	Hillshade_3pm	\
0	221.0	232.0	148.0	
1	220.0	235.0	151.0	
2	234.0	238.0	135.0	
3	238.0	238.0	122.0	
4	220.0	234.0	150.0	

	Horizontal_Distance_To_Fire_Points	...	Soil_Type32	Soil_Type33	\
0	6279.0	...	0.0	0.0	
1	6225.0	...	0.0	0.0	
2	6121.0	...	0.0	0.0	
3	6211.0	...	0.0	0.0	
4	6172.0	...	0.0	0.0	

	Soil_Type34	Soil_Type35	Soil_Type36	Soil_Type37	Soil_Type38	\
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	

	Soil_Type39	Soil_Type40	Cover_Type
0	0.0	0.0	5
1	0.0	0.0	5
2	0.0	0.0	2
3	0.0	0.0	2
4	0.0	0.0	5

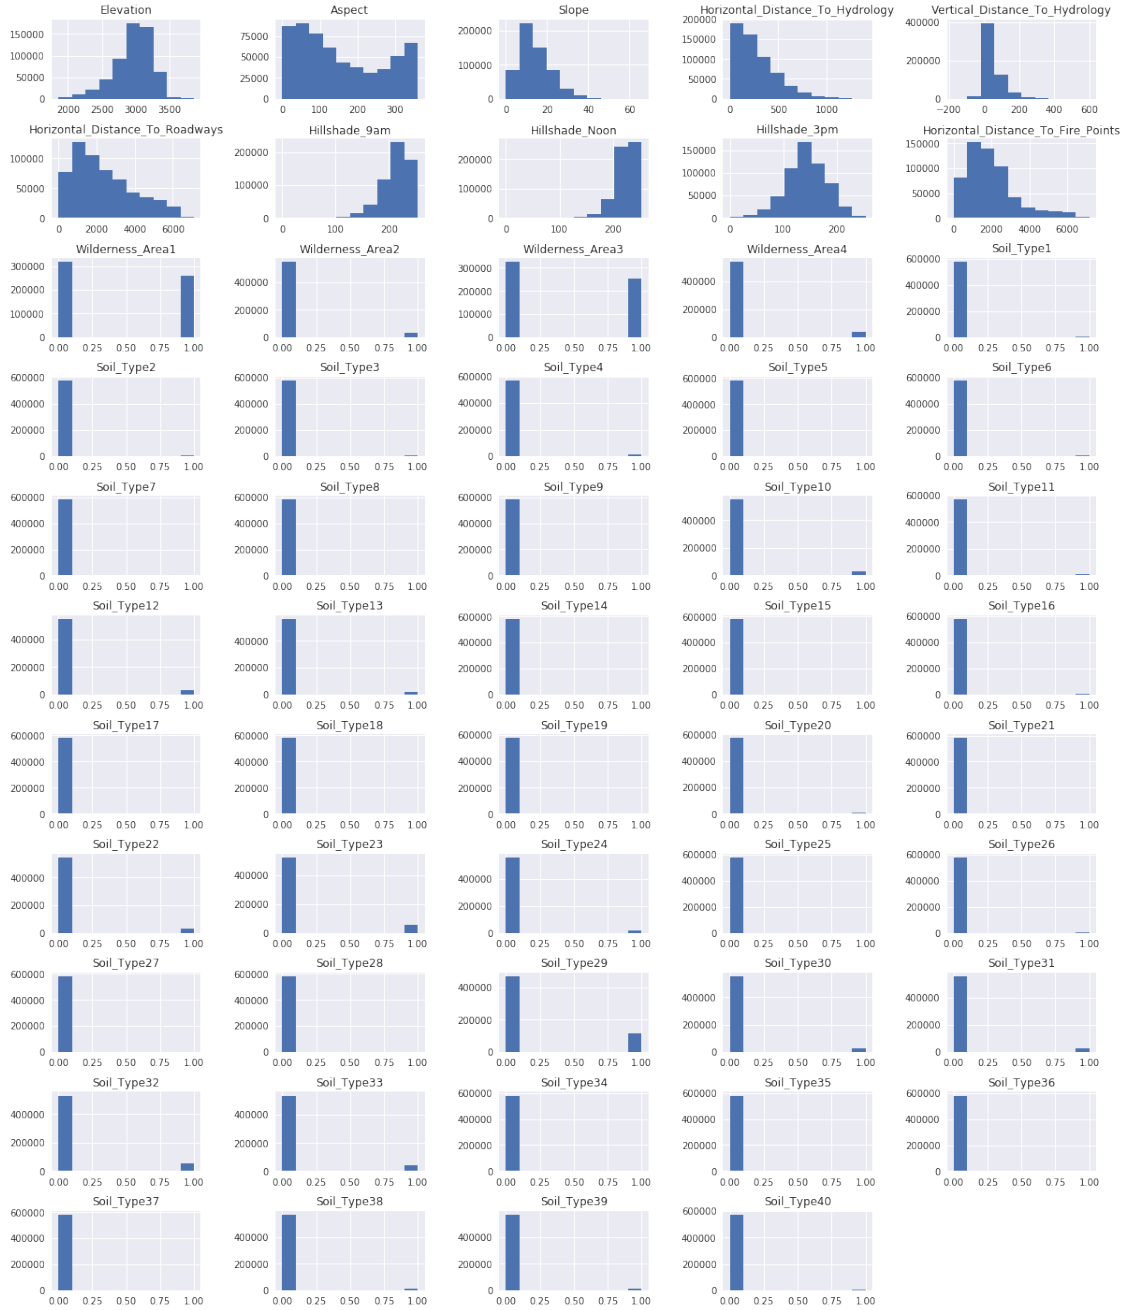
[5 rows x 55 columns]

```
In [11]: cov.shape
```


```
Out[11]: (581012, 55)
```

0.0.1 Task 2.1: Visualize the univariate distribution of each feature, and the distribution of the target

```
In [12]: fig = plt.figure(figsize=(20,25))
fig.subplots_adjust(hspace=0.5, wspace=0.5)
for i in range(1, 55):
    ax = fig.add_subplot(11, 5, i)
    cov[cov.columns[i-1]].hist(ax=ax)
    ax.set(title=cov.columns[i-1])
```



```
In [13]: #Distribution of target
plt.hist(cov['Cover_Type'],range=(0.5,7.5),bins=7)
_ = plt.title('Cover Type Histogram')
```



0.0.2 Task 2.2: Split data into training and test set. Evaluate Logistic Regression, linear support vector machines and nearest centroids using cross-validation

```
In [14]: X_train, X_test, y_train, y_test = train_test_split(data['data'], data['target'], random_
```

```
        X_train = X_train[:50000] y_train = y_train[:50000]
```

```
In [15]: #Evaluate Logistic Regression
        logit_score = np.mean(cross_val_score(LogisticRegression(tol=0.0001), X_train, y_train))
```

```
In [16]: #Linear SVM
        linearSVM_score = np.mean(cross_val_score(LinearSVC(tol=0.0001, dual=False), X_train, y
```

```
In [17]: #Nearest Centroid
        NC_score = np.mean(cross_val_score(NearestCentroid(), X_train, y_train))
```

```
In [18]: print(logit_score)
        print(linearSVM_score)
        print(NC_score)
```

```
0.708256150789
```

```
0.705947542451
```

```
0.194639694306
```

```

In [19]: #Use standardScaler
        scaler = StandardScaler()
        X_train_scaled = scaler.fit_transform(X_train)

In [20]: slogit_score = np.mean(cross_val_score(LogisticRegression(tol=0.0001), X_train_scaled,
        slinearSVM_score = np.mean(cross_val_score(LinearSVC(tol=0.0001, dual=False), X_train_scaled,
        sNC_score = np.mean(cross_val_score(NearestCentroid(), X_train_scaled, y_train))

In [21]: print(slogit_score)
        print(slinearSVM_score)
        print(sNC_score)

0.715418375449
0.712604890841
0.548706947943

```

Logistic regression and Linear SVM have close scores - Logistic regression is slightly better, while Nearest Centroid has very bad score compared to other 2 algorithms. When applying StandardScaler, all 3 algorithms improve in score: logistic regression from 70.8% to 71.5%, linear SVM from 70.6% to 71.3%. Nearest Centroid improves the greatest from 19.5% to 54.9%, although final score is still far behind other 2 algorithms.

0.0.3 Task 2.3: Tune the parameters using GridSearchCV

```

In [22]: #Tune logistic regression
        param_grid_logit = {'C':np.logspace(-2,2,6)}
        grid_logit = GridSearchCV(LogisticRegression(tol=0.0001,random_state=20), param_grid_logit)
        grid_logit.fit(X_train_scaled, y_train)
        print(grid_logit.best_params_)
        print(grid_logit.best_score_)

{'C': 100.0}
0.715496409713

In [23]: #Tune Linear SVM
        param_grid_svm = {'C':np.logspace(-2,2,6)}
        grid_svm = GridSearchCV(LinearSVC(tol=0.0001, dual=False, random_state=20), param_grid_svm)
        grid_svm.fit(X_train_scaled, y_train)
        print(grid_svm.best_params_)
        print(grid_svm.best_score_)

{'C': 2.5118864315095824}
0.712616377401

```

```

In [24]: #Tune Nearest Centroid
        param_grid_NC = {'shrink_threshold':np.linspace(0,20,11)}

```

```

grid_NC = GridSearchCV(NearestCentroid(), param_grid_NC)
grid_NC.fit(X_train_scaled, y_train)
print(grid_NC.best_params_)
print(grid_NC.best_score_)

```

```

{'shrink_threshold': 0.0}
0.54870696876


```

In [25]: *#Visualize GridSearch results*

```

plt.semilogx(param_grid_logit['C'],grid_logit.cv_results_['mean_train_score'],label='Train accuracy')
plt.semilogx(param_grid_logit['C'],grid_logit.cv_results_['mean_test_score'],label='Test accuracy')
_ = plt.title('Logistic Regression C parameter')
_ = plt.legend()


```




```

In [26]: plt.semilogx(param_grid_svm['C'],grid_svm.cv_results_['mean_train_score'],label='Train accuracy')
plt.semilogx(param_grid_svm['C'],grid_svm.cv_results_['mean_test_score'],label='Test accuracy')
_ = plt.title('Linear SVM C parameter')
_ = plt.legend()

```



```
In [27]: plt.plot(param_grid_NC['shrink_threshold'],grid_NC.cv_results_['mean_train_score'],label='mean_train_score')
plt.plot(param_grid_NC['shrink_threshold'],grid_NC.cv_results_['mean_test_score'],label='mean_test_score')
_ = plt.title('Nearest Centroid Shrink Threshold')
_ = plt.legend()
```



After GridSearch, the scores for all 3 algorithms do not improve by too much.

0.0.4 Task 2.4: Change the cross-validation strategy from 'stratified k-fold' to 'kfold' with shuffling

```
In [28]: #Change cross-validation strategy to KFold with shuffling and run GridSearch
        kfoldshuffle = KFold(shuffle=True, random_state=1)
```

```
In [29]: #Tune logistic regression with KFold with shuffling
        grid_logit_kfold = GridSearchCV(LogisticRegression(tol=0.0001,random_state=20), param_g
        grid_logit_kfold.fit(X_train_scaled, y_train)
        print(grid_logit_kfold.best_params_)
        print(grid_logit_kfold.best_score_)
```

```
{'C': 100.0}
0.715503294252
```

```
In [30]: #Tune Linear SVM with KFold with shuffling
        grid_svm_kfold = GridSearchCV(LinearSVC(tol=0.0001, dual= False, random_state=20), para
        grid_svm_kfold.fit(X_train_scaled, y_train)
        print(grid_svm_kfold.best_params_)
        print(grid_svm_kfold.best_score_)
```

```
{'C': 2.5118864315095824}
0.712731119724
```

```
In [31]: #Tune Nearest Centroid with KFold with shuffling
        grid_NC_kfold = GridSearchCV(NearestCentroid(), param_grid_NC,cv=kfoldshuffle)
        grid_NC_kfold.fit(X_train_scaled, y_train)
        print(grid_NC_kfold.best_params_)
        print(grid_NC_kfold.best_score_)
```

```
{'shrink_threshold': 0.0}
0.548782698694
```

```
In [32]: #Change random seed of the shuffling
        kfoldshuffle2 = KFold(shuffle=True, random_state=2)
```

```
In [33]: #Tune logistic regression with KFold with shuffling
        grid_logit_kfold2 = GridSearchCV(LogisticRegression(tol=0.0001,random_state=20), param_
        grid_logit_kfold2.fit(X_train_scaled, y_train)
        print(grid_logit_kfold2.best_params_)
        print(grid_logit_kfold2.best_score_)
```

```
{'C': 15.848931924611142}  
0.715455102476
```

```
In [34]: #Tune Linear SVM with KFold with shuffling  
grid_svm_kfold2 = GridSearchCV(LinearSVC(tol=0.0001,dual=False, random_state=20), param_g  
grid_svm_kfold2.fit(X_train_scaled, y_train)  
print(grid_svm_kfold2.best_params_)  
print(grid_svm_kfold2.best_score_)
```

```
{'C': 15.848931924611142}  
0.71274718365
```

```
In [35]: #Tune Nearest Centroid with KFold with shuffling  
grid_NC_kfold2 = GridSearchCV(NearestCentroid(), param_grid_NC,cv=kfoldshuffle2)  
grid_NC_kfold2.fit(X_train_scaled, y_train)  
print(grid_NC_kfold2.best_params_)  
print(grid_NC_kfold2.best_score_)
```

```
{'shrink_threshold': 0.0}  
0.548617469748
```

```
In [52]: #Change random state of split and run GridSearch  
X_train3, X_test3, y_train3, y_test3 = train_test_split(data['data'],data['target'],ran  
X_train3_scaled = scaler.fit_transform(X_train3)
```

```
In [56]: #Tune logistic regression with different random split  
grid_logit_split2 = GridSearchCV(LogisticRegression(tol=0.0001,random_state=3), param_g  
grid_logit_split2.fit(X_train3_scaled, y_train3)  
print(grid_logit_split2.best_params_)  
print(grid_logit_split2.best_score_)
```

```
{'C': 100.0}  
0.714791891848
```

```
In [62]: #Tune Linear SVM with with different random split  
grid_svm_split2 = GridSearchCV(LinearSVC(tol=0.0001,dual=False,random_state=3), param_g  
grid_svm_split2.fit(X_train3_scaled, y_train3)  
print(grid_svm_split2.best_params_)  
print(grid_svm_split2.best_score_)
```

```
{'C': 2.5118864315095824}  
0.712142942554
```

```
In [58]: #Tune Nearest Centroid with different random split
grid_NC_split2 = GridSearchCV(NearestCentroid(), param_grid_NC)
grid_NC_split2.fit(X_train3_scaled, y_train3)
print(grid_NC_split2.best_params_)
print(grid_NC_split2.best_score_)

{'shrink_threshold': 0.0}
0.548819416237
```

The parameters after GridSearch do not change after changing from StratifiedKFold to KFold with shuffling. However the parameters do change when we change the random state of the shuffling, from (100,2.5) for Logistic Regression and Linear SVM to (15.8,15.8). When changing the random state of the split into training and test data, the parameters also do not change compared to original set of parameters. In the 3 cases, parameters of Nearest Centroid do not change.

0.0.5 Task 2.5: Visualize the coefficients for LogisticRegression and Linear Support Vector Machines

```
In [42]: #Logistic regression model
logitmodel = LogisticRegression(tol=0.0001,C=grid_logit.best_params_['C'],random_state=
logitmodel.fit(X_train_scaled,y_train)
#Linear SVM model
linearSVMmodel = LinearSVC(tol=0.0001, dual=False, C=grid_svm.best_params_['C'], random
linearSVMmodel.fit(X_train_scaled, y_train)

Out[42]: LinearSVC(C=2.5118864315095824, class_weight=None, dual=False,
fit_intercept=True, intercept_scaling=1, loss='squared_hinge',
max_iter=1000, multi_class='ovr', penalty='l2', random_state=20,
tol=0.0001, verbose=0)

In [43]: fig = plt.figure(figsize=(15,15))
for i in range(1,8):
    ax = fig.add_subplot(4,2,i)
    ax.scatter(range(X_train.shape[1]), logitmodel.coef_[i-1],label='Logistic Regression')
    ax.scatter(range(X_train.shape[1]), linearSVMmodel.coef_[i-1],label='Linear SVM')
    ax.set(title='Cover Type '+str(i))
plt.legend()
```

